



# *Workshop 9 – Serverless (Function as a Service (FaaS))*

Luca Morandini

Cloud Architect – Melbourne eResearch Group

University of Melbourne

[luca.morandini@unimelb.edu.au](mailto:luca.morandini@unimelb.edu.au)

# Outline of the Lecture

- **Part 1: What is FaaS**
  - Serverless / FaaS defined
  - Why functions?
  - FaaS services and frameworks
- **Part 2: More About Functions**
  - Functions with side-effects
  - Stateful/stateless functions
  - Synchronous/Asynchronous functions
- **Part 3: Fn Workshop**
  - Introduction to Fn
  - Installation
  - Function Deployment

## **Part 1: What is Function as a Service (FaaS) ?**

# Disambiguation

- FaaS is also known as *Serverless computing* (more catchy, but less precise)
- The idea behind Serverless/FaaS is to develop software applications **without bothering with the infrastructure** (especially **scaling-up and down** as load increases or decreases): the service provider does it for you
- Therefore, it is more *Server-unseen than Server-less*
- A FaaS service allows functions to be added, removed, updated, executed, and auto-scaled
- FaaS is, in a way, an extreme form of *microservice architecture*

# Why Functions?

- A *function* in computer science is a piece of code that takes in parameters and returns a value
- Functions are the founding concept of *functional programming* -one of the oldest programming paradigms
- Functions are, or should be, free of *side-effects*, *ephemeral*, and *stateless*, which make them ideal for parallel execution and rapid scaling-up and -down, hence their use in FaaS (more on this topic later)

# Why FaaS?

- Simpler deployment (the service provider takes care of the infrastructure)
- Reduced computing costs (only the time during which functions are executed is billed)
- Reduced application complexity due to loosely-coupled architecture

# FaaS Applications?

- Functions are triggered by events
- Functions can call each other
- Functions and events can be combined to build software applications
- For instance: a function can be triggered every hour (say, to compress log files), or every time disk space on a volume is scarce (to remove old log files), or when a pull-request is merged in GitHub, or when a message is stored in a queue
- Combining event-driven scenarios and functions resembles how User Interface software is built: user actions trigger the execution of pieces of code

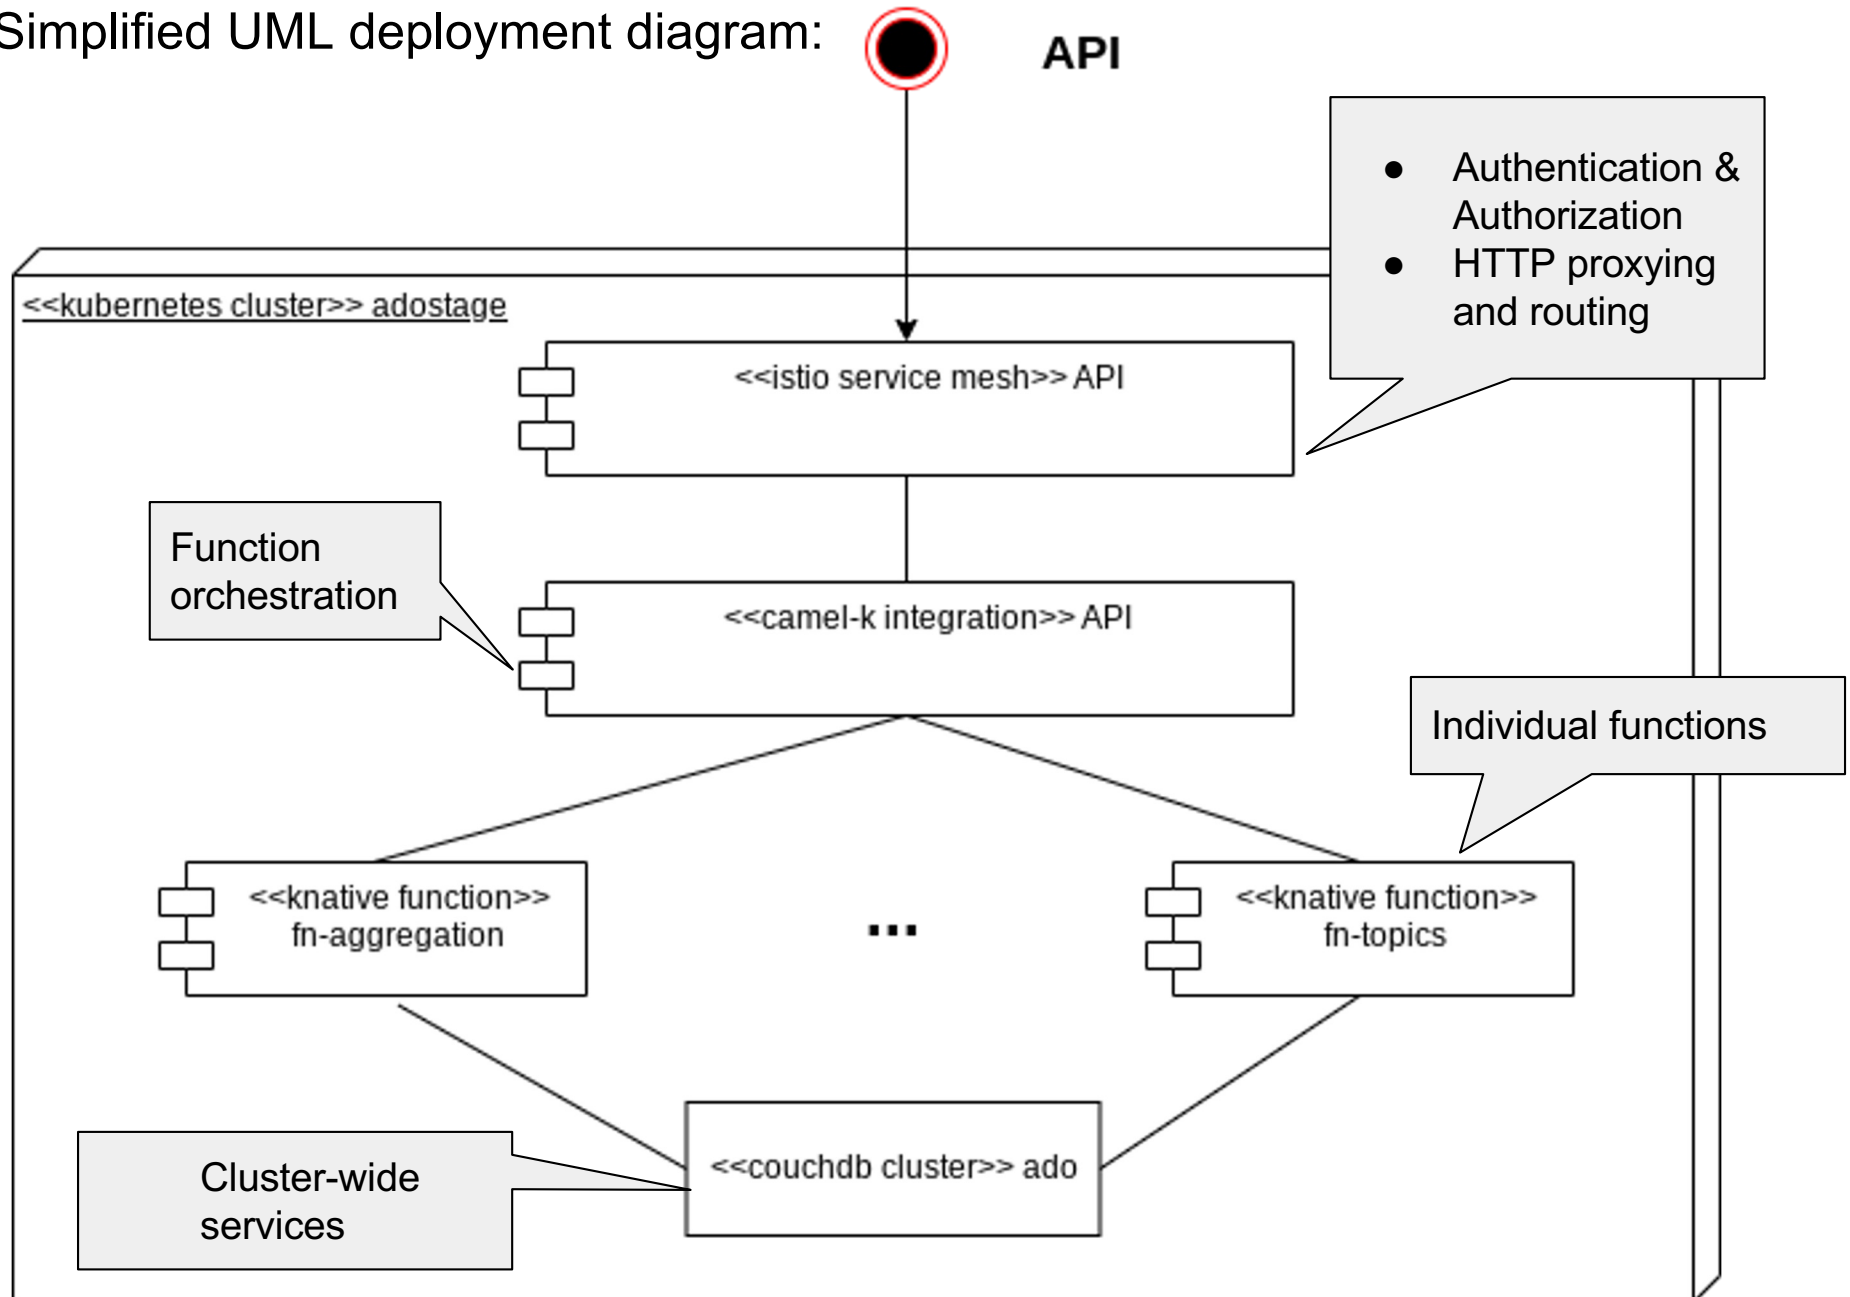
# Serverless Application: ADO #1

- The ADO backend is build on a Kubernetes cluster. Kubernetes is a container orchestration platform that allows to distribute containers across a number of VMs
- ADO use a few frameworks:
  - **Knative**: serverless framework for deploying and managing the life-cycle of functions including auto-scaling
  - **Camel-K**: orchestration of Knative functions, such as function composition (flow), caching of function results, message queues
  - **Istio**: a service mesh that takes care of authentication and authorization, HTTP header management, routing



# Serverless Application: ADO #2

- Simplified UML deployment diagram:



# Serverless Application: ADO #3

- The API is defined as an Open API specification:

```
/data/collections/{collection}/summary:
  get:
    summary: Return information about a collection
    operationId: analysisAggregateSummary
    parameters:
      - name: collection
        in: path
        description: collection name
        required: true
```

- A Camel-K integration routes the `operationID` to actual Knative functions:

```
from("direct:analysisAggregateSummary")
  .to("knative:endpoint/fn-aggregation")
```

- Functions are Kubernetes pods (sets of related containers):

```
fn-aggregation-00001-deployment-cdff847d-mnj6n      2/2
Running      0      20d
...
fn-textsearch-00001-deployment-6c6c589f69-hjqxl      2/2
Running      0      20d
fn-topics-00001-deployment-6fdbcb7b556-9wkpp         2/2
Running      0      20d
```

- Functions are independent of each other, the API specification is written in YAML, and the Camel-K integration routing could be written in XML/Java/JavaScript/etc. Therefore a developer team can mix and match different languages, libraries, and even operating systems without worrying about the deployment on the actual computing resources

# FaaS Services and Frameworks

- The first widely available FaaS service was Amazon's AWS Lambda. Since then Google Cloud Functions (part of Firebase) and Azure Functions by Microsoft
- All of the FaaS above **allow functions to use the services of their respective platforms**, thus providing a rich development environment
- There are several open-source frameworks (*funtainers* - or *functions containers*) such as *Apache OpenWhisk*, *OpenFaas*, *Fn*, and *Knative*
- The main difference between proprietary FaaS services and open-source FaaS frameworks is that the latter can be deployed on your cluster, peered into, disassembled, and improved by you.

## **Part 2: More About Functions**

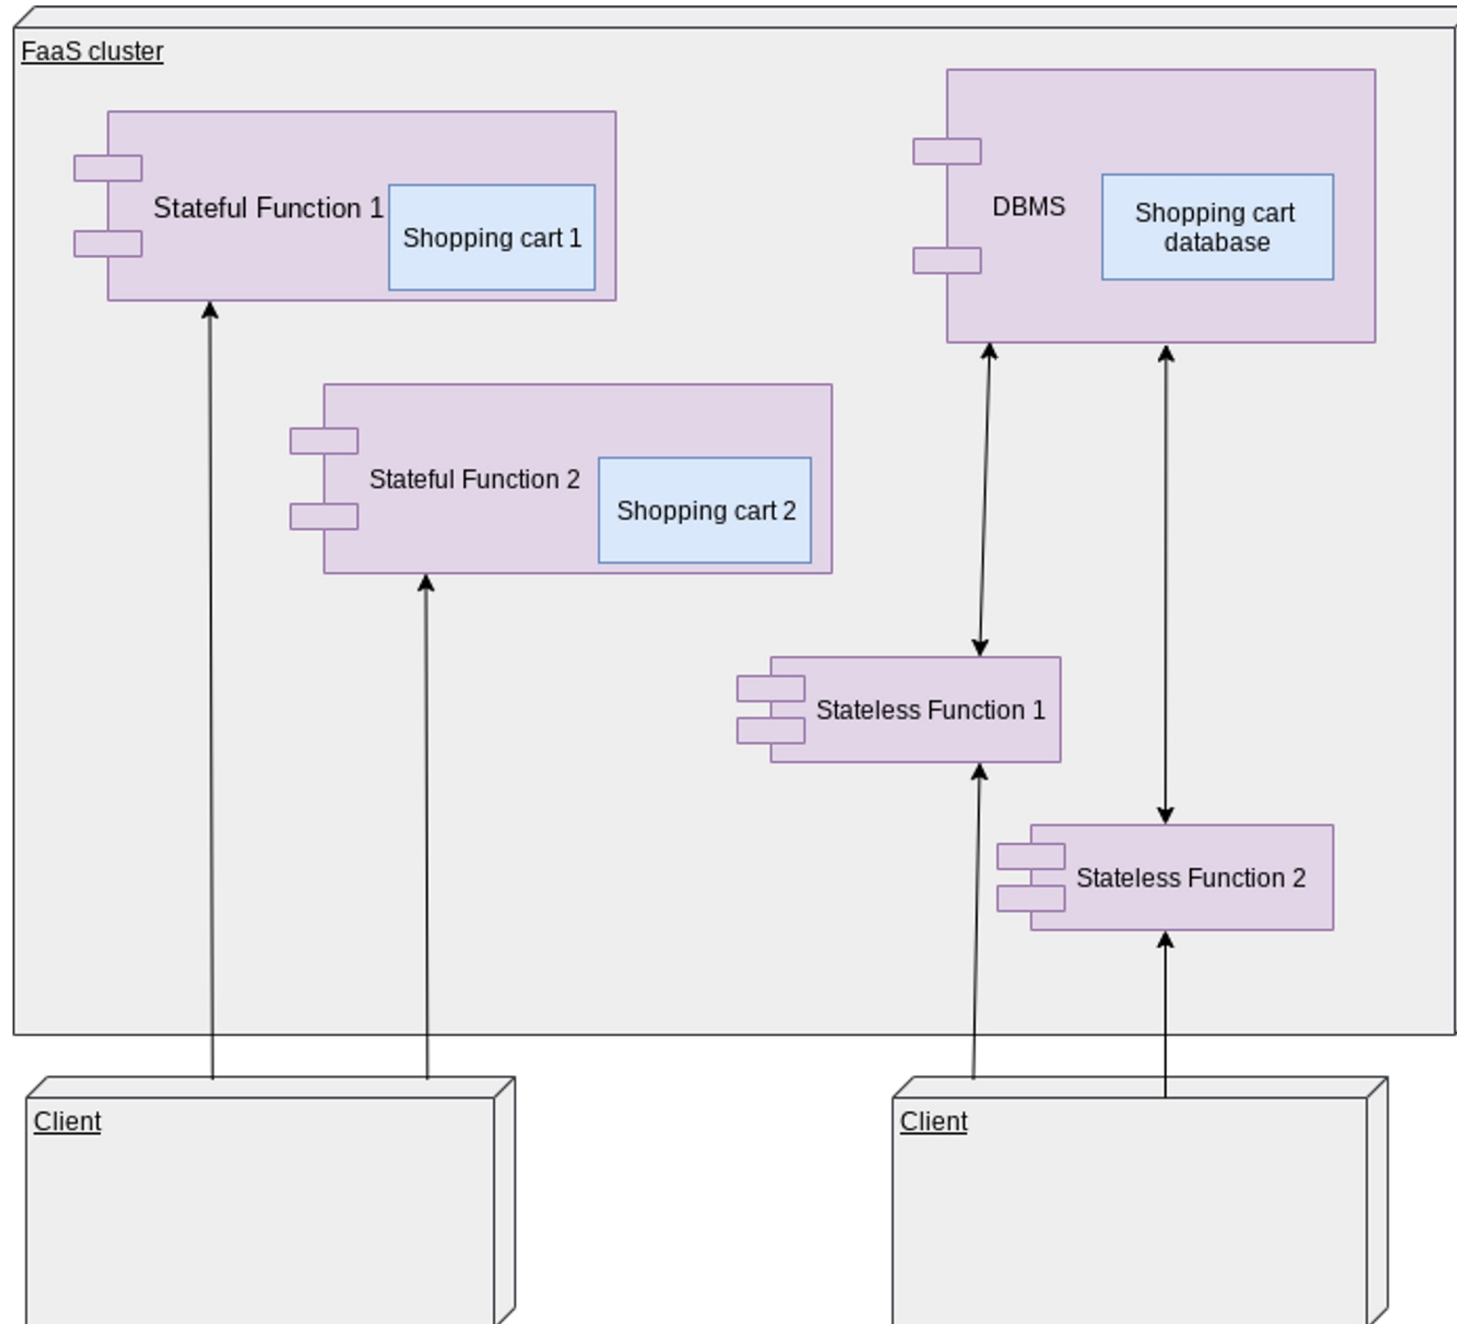
# Side-effect Free Functions

- A function that does not modify the state of the system is said to be *side-effect free* (for instance, a function that takes an image and returns a thumbnail of that image)
- A function that changes the system somehow is not side-effect free (for instance, a function that writes to the file system the thumbnail of an image)
- Side-effect free functions can be run in parallel, and are guaranteed to return the same output given the same input
- Side-effects, however, are almost inevitable in a relatively complex system. Therefore consideration must be given on how to make functions with side effects run in parallel, as typically required in FaaS environments.

# Stateful/Stateless Functions #1

- A subset of functions with side-effects is composed of *stateful functions*
- A **stateful** function is one whose output changes in relation to internally stored information (**hence its input cannot entirely predict its output**), e.g. a function that adds items to a “shopping cart” and retains that information internally
- Conversely, a *stateless function* is one that does not store information internally, e.g. adding an item to a “shopping cart” stored in a DBMS service and not internally would make the function above stateless, but not side-effect free.
- This is important in FaaS services since there are multiple instances of the same function, and there is no guarantee the same user would call the same function instance twice

# Stateful/Stateless Functions #2



# Synchronous/Asynchronous Functions

- By default functions in FaaS are *synchronous*, hence they return their result immediately (or almost so)
- However, there may be functions that take longer to return a result, hence they incur timeouts and lock connections with clients in the process, hence it is better to transform them into *asynchronous functions*
- Asynchronous functions return a code that informs the client that the execution has started, and then trigger an event when the execution completes
- In more complex cases a *publish/subscribe pattern* involving a queuing system can be used to deal with asynchronous functions



## **Part 3: Fn Workshop**

# Introduction to Fn

- Fn is an open-source framework that uses **Docker containers** to deliver FaaS functionality
- **Every function in Fn is a Docker container**, ensuring **loose coupling** between functions (functions can be written in different languages and mixed freely)
- By using Docker containers as functions, Fn allow to **freely mix different languages and environments** at the cost of decreased performance, as containers are inherently heavier than threads. However, by using a bit of finesse, a container with a single executable, can weight only a few MBs (check out [https://hub.docker.com/\\_/scratch](https://hub.docker.com/_/scratch))

# ...More on Fn

- Fn is the technology behind Oracle Functions (the serverless service of Oracle Cloud)
- Fn could be deployed on Kubernetes to manage cluster of nodes on which functions are run
- Fn allows both synchronous and asynchronous functions
- With Fn Flow, functions can be composed efficiently
- Fn can add more Docker containers when a function is called more often, and remove containers when the function is called less often

# Calling Functions Defined in Fn

- Calling a function in Fn can be done via **POST** HTTP request, as in:  

```
curl -XPOST\  
  "http://0.0.0.0:8080/t/wcapp/wcmp"\  
  --data "Lorem ipsum dixit"
```
- Functions can be grouped into an “App”, so that they are arranged into hierarchical URLs (i.e. `/myapp/myfunction`). By default the paths follow the directory structure of the source code
- Function can be triggered by HTTP requests, but also by other classes of events, such as new messages in message queues (publish/subscribe)

# Node.js Function in Fn

```
const fdk = require ('@fnproject/fdk');

fdk.handle ( (input) => {
  let counts = {};
  input.toLowerCase ()
    .split(/\W+/)
    .filter ((w) => {
      return w.length > 1;
    })
    .forEach ((w) => {
      counts[w] = (counts[w] ? counts[w] + 1 : 1);
    });
  return counts;
})
```

# Python Function in OpenFaaS

```
import io
import json
import logging
from fdk import response

def handler(ctx, data: io.BytesIO = None):
    try:
        counts = dict()
        logger = logging.getLogger()
        for word in data.getvalue().split():
            counts[word.decode('utf-8')] =
                (counts.get(word.decode('utf-8')) or 0) + 1
    except (Exception, ValueError) as ex:
        logging.getLogger().info('error: ' + str(ex))
    return response.Response(
        ctx, response_data=json.dumps(counts),
        headers={"Content-Type": "application/json"}
    )
```

- “ctx” is a Context object than contains metadata about the function and its invocation: additional configuration parameters, the HTTP method, HTTP headers, etc.

# Workshop repository

- Use `git` to clone the repository at:  
<https://gitlab.unimelb.edu.au/feit-comp90024/comp90024>
- Go to the `fn` directory
- Follow the instructions in the `README.md`